

Unit 4 - Repeating and finishing a project

Draft

Clara Himmelbauer

General overview

What is Quarto?

Quarto enables you to write code and text “as one”. It is the Posit (the company behind RStudio) version of Jupyter Notebooks, which became popular for python in recent years.

How does it work?

Quarto has a source and a visual editor. We will be mostly working in the visual editor. If you go directly to the source, you will see that the underlying “language” is Markdown - a typography language used a lot in programming and text designing. You can do everything with the buttons in the menu above, but you might still learn the most basic markdown shortcuts, e.g. for inserting code or making headings.

To render your document, click the **Render** button above. You might want to change your settings to render inside the RStudio viewer beforehand.

Resources

There are numerous markdown guides and cheatsheets available, for example [this one](#). The general guide for Quarto can be found [here](#) on quarto.org.

Preparation

YAML code at the top

The Code chunk at the top is called “YAML”, it specifies some of the metadata used throughout the document. Here, you can also change the type of your document output. For working, HTML is suggested. You can always change it to pdf later when exporting the document.

Packages

At first, we need to attach packages. Besides the usual `tidyverse` package, we also need `rmarkdown` for this.

Another thing: Whenever you start a code chunk, you can give it some options using `#|`. An overview of execution options can be found [here](#).

```
# install.packages("rmarkdown")
library("rmarkdown")
library("tidyverse")
library("ggplot2")
library("stargazer")
```

Data

Next, we need to import the data. We are going to use the same data as previous weeks - the SILC-Data.

Before we do that we must however take care where our `qmd` file is located and saved, using the `getwd()` command. Then we have multiple options to deal with that. The easiest one is to re-define the path we are working at. Make sure, that you change this path whenever you are working on another PC.

```
getwd()
```

```
[1] "C:/Users/chimmelb/OneDrive - WU Wien/Dokumente/Flinta-R-Tut-Summer25/analysis"
```

```
parent_path <- c("C:/Users/chimmelb/OneDrive - WU Wien/Dokumente/Flinta-R-Tut-Summer25")
hh <- readRDS(paste0(parent_path, "/", "data/silc_hh_new.RDS"))
ind <- readRDS(paste0(parent_path, "/", "data/silc_indiv_new.RDS"))
all <- readRDS(paste0(parent_path, "/", "data/silc_all.RDS"))
```

We might also want to look at our data, so let's do this. `glimpse` and `str` both are nice ways to look at data.

```
glimpse(hh)
```

Rows: 116,569

Columns: 9

```
$ year      <int> 2008, 2008, 2008, 2008, 2008, 2008, 2008, 2008, 2008, 2008, 20~
$ country   <chr> "AT", "AT", "AT", "AT", "AT", "AT", "AT", "AT", "AT", "AT", "A~
$ hid       <int> 4, 5, 6, 7, 9, 10, 11, 23, 12, 13, 16, 17, 18, 19, 20, 24, 25,~
$ hinc      <dbl> 16792.40, 58727.60, 13044.20, 19858.20, 74073.52, 47715.87, 24~
$ hsize     <int> 1, 5, 1, 2, 4, 3, 2, 3, 4, 3, 2, 4, 3, 4, 2, 1, 1, 5, 1, 5, 2,~
$ heqinc    <dbl> 16792.40, 24469.83, 13044.20, 13238.80, 29629.41, 23857.94, 16~
$ region    <chr> "AT3", "AT1", "AT3", "AT3", "AT1", "AT1", "AT2", "AT3", "AT1",~
$ hweight   <dbl> 845.0613, 199.2002, 1060.7000, 488.6925, 1155.0970, 550.5005, ~
$ degurba   <int> 2, 3, 1, 2, 3, 3, 3, 1, 2, 3, 1, 1, 2, 1, 1, 2, 2, 2, 3, 2, 1,~
```

```
str(ind)
```

```
'data.frame':  229131 obs. of  12 variables:
 $ year      : int  2008 2008 2008 2008 2008 2008 2008 2008 2008 2008 2008 ...
 $ country   : chr  "AT" "AT" "AT" "AT" ...
 $ pid       : int  1001 100101 1002 100201 100202 1003 100401 100402 100501 100502 ...
 $ pweight   : num  551 583 551 1103 1103 ...
 $ birthyear : int  1927 1977 1929 1966 1970 1957 1942 1947 1927 1932 ...
 $ sex       : int  1 1 2 1 2 2 1 2 1 2 ...
 $ educ      : int  2 5 2 1 3 2 5 3 2 2 ...
 $ empstatus : int  NA 2 NA 3 NA 3 3 3 2 NA ...
 $ workinghours: int  NA 42 NA 32 NA 40 NA NA NA NA ...
 $ health    : int  3 1 2 2 1 2 3 3 3 3 ...
 $ gross_income: num  0 5037 0 7989 0 ...
 $ hid       : int  10 1001 10 1002 1002 10 1004 1004 1005 1005 ...
```

Recap: dplyr and ggplot

We're going to make a simple plot: In 2008, mean equivalized household income by country and degree of urbanization.

Data wrangling

```
df <- hh %>%
  filter(year == 2008) %>%
  group_by(country, degurba) %>%
  summarise(heqinc = weighted.mean(heqinc, hweight))
```

Table

Let's render the table first.

There are numerous ways to make nice tables. For presentation you might look in to `kable()`, `kableExtra`, and `gt`. Furthermore, you might want to export tables to LaTeX. Here, `stargazer` is the typical option. Don't forget to install and attach the packages above though.

```
knitr::kable(df)
```

Table 1: Mean equivalized household income by country and Degree of Urbanization.

country	degurba	heqinc
AT	1	21893.91
AT	2	22413.47
AT	3	20394.14
DE	1	20968.12
DE	2	20910.77
DE	3	17842.63
IT	1	18798.08
IT	2	17320.09
IT	3	15671.60

Plot

Now we make a nice barchart. Before we do that, however, let's make an internal link: The data used for the following plot is already displayed in Table ??.

Also, our `degurba` variable is still a numeric (integer). So let's recode it to a factor.

In the following code chunk, I use the option `#| echo: false`. It prevents the R-code from being shown and in the resulting document, only the ggplot is included. If you want this settings to apply to the whole document, you can include it at the YAML code chunk at the top of the quarto document.

Practice time

So now it's time for you to practice. Make a plot, where you have income on the y axis, and employment hours on the x axis. There should be 3 different lines, one for each country. The year we are doing this for is 2013.




Figure 1: Mean equivalized household income by country and DEGURBA.

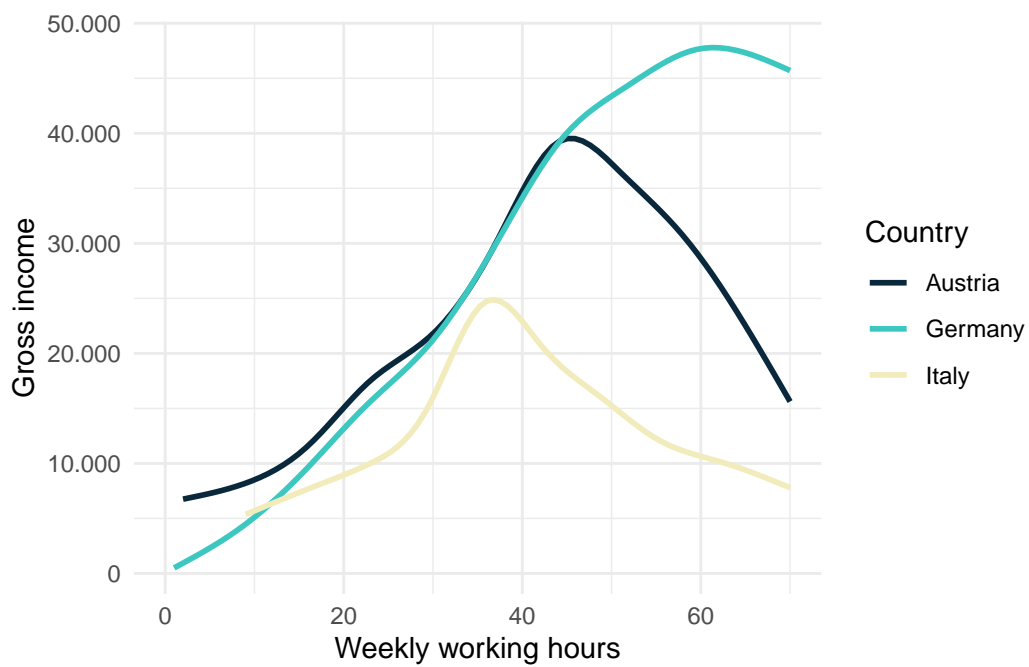


Figure 2: Income by employment hours, 2013.